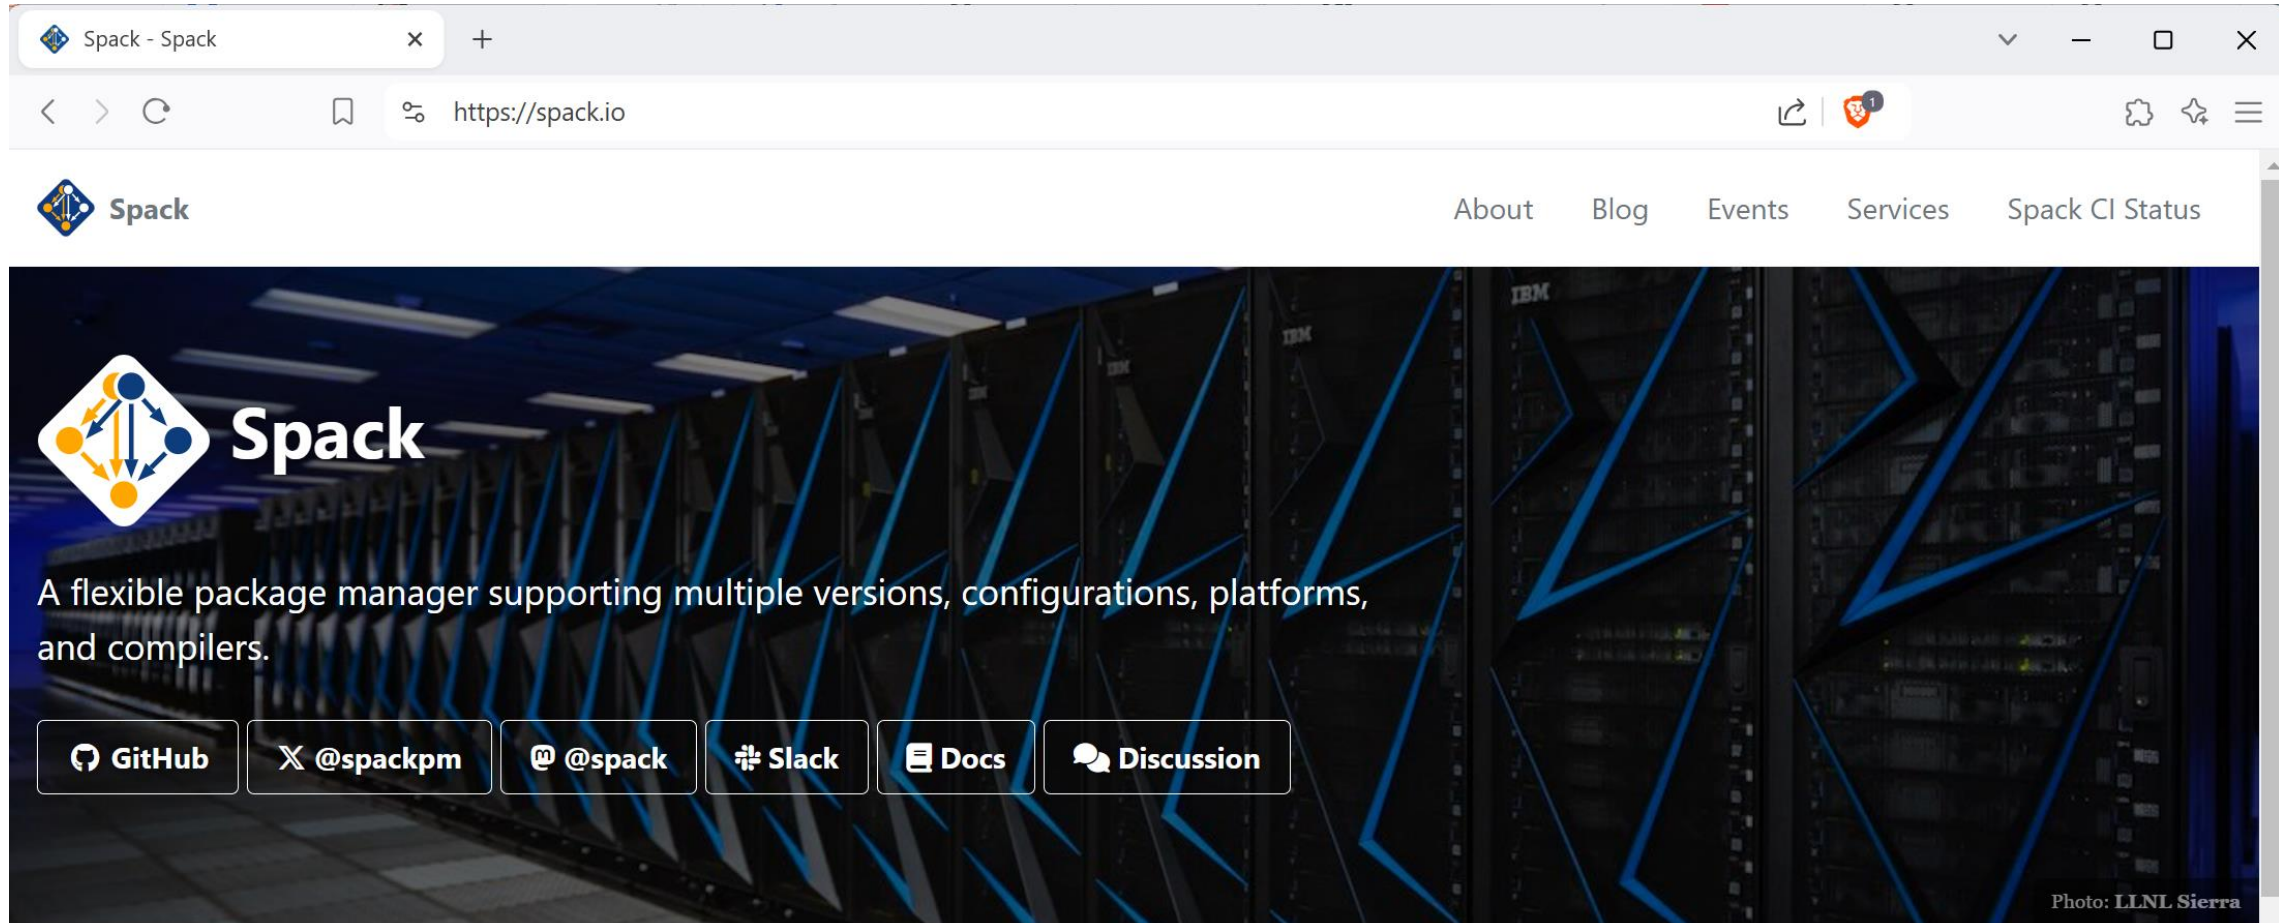




# Spack

A package manager for HPC systems  
Peter Larsson - PDC

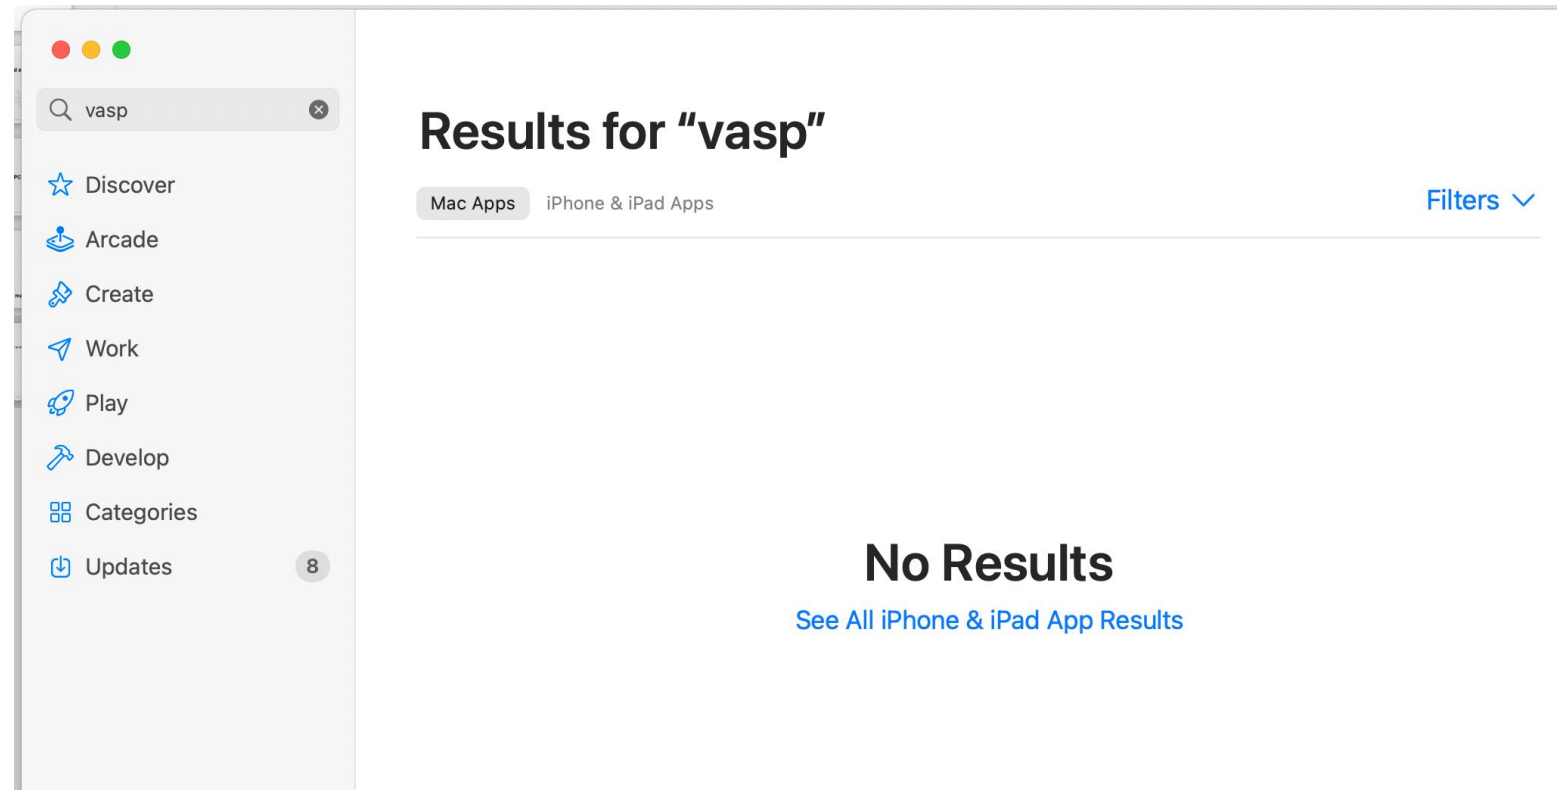


The screenshot shows a web browser window with the Spack website. The browser's address bar shows 'https://spack.io'. The website has a navigation bar with links for 'About', 'Blog', 'Events', 'Services', and 'Spack CI Status'. The main content area features a large background image of server racks with blue geometric overlays. On the left, there is a Spack logo (a diamond shape with four colored circles and arrows) and the word 'Spack' in a large, white, sans-serif font. Below the logo, the text reads: 'A flexible package manager supporting multiple versions, configurations, platforms, and compilers.' At the bottom of this section, there are six buttons with icons and text: 'GitHub', '@spackpm', '@spack', 'Slack', 'Docs', and 'Discussion'. In the bottom right corner of the main image, there is a small credit: 'Photo: LLNL Sierra'.

## Welcome to Spack!

Spack is a package manager for [supercomputers](#), Linux, and macOS. It makes installing scientific software easy. Spack isn't tied to a particular language; you can build a software stack in [Python](#) or R, link to libraries written in C, C++, or Fortran, and easily [swap compilers](#) or target [specific microarchitectures](#). Learn more [here](#).

# What if we had an App Store for the HPC cluster?



*(VASP is used to simulate materials at the atomic level.  
It is probably the most popular HPC application?)*

# Phone: App store

Consider the modern way of installing software on a new phone:

- You log in to the app store...
- your previously installed software from the other phone can/will be downloaded automatically...
- ...and **optimized for the new hardware automatically** (it might be a completely new CPU/GPU etc)
- All Apps are sandboxed and independent of each other.
- New Apps can be added by searching for their name, and just click install.

# HPC cluster: Spack

Spack is the closest thing to an "App store" model for HPC:

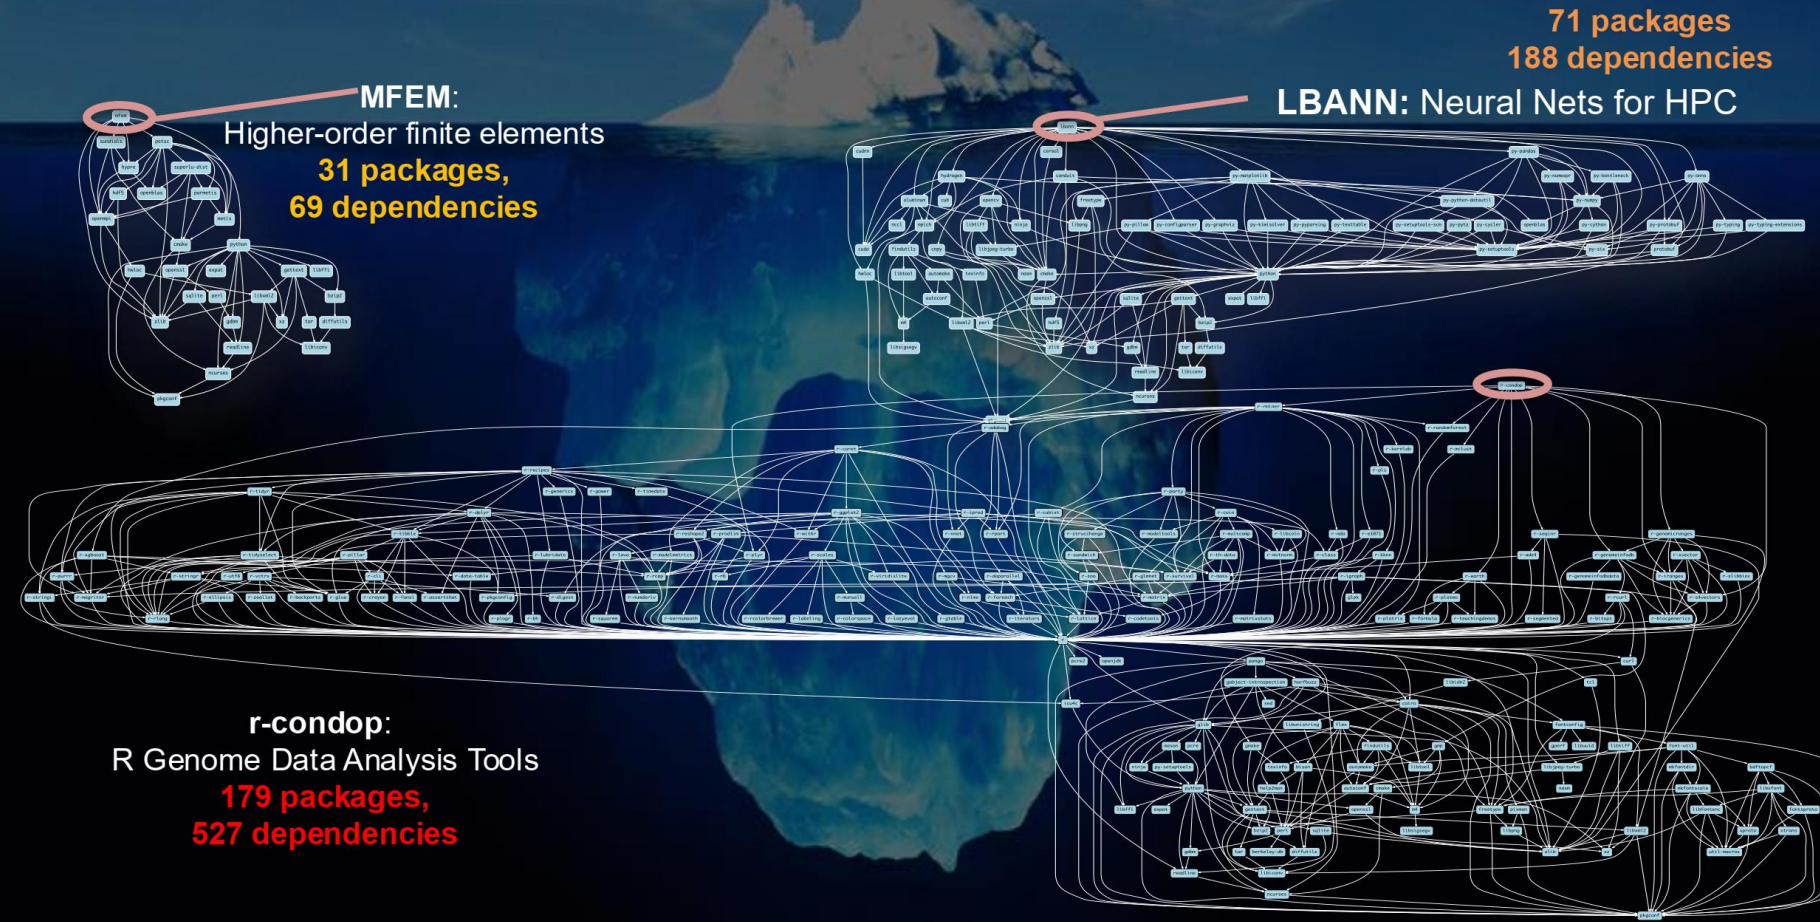
- Install Spack
- Copy your `spack.yaml` file from your other HPC cluster. It defines your "software environment".
- Recreate the environment on the new machine  
`spack env create $NAME spack.yaml`
- Tell Spack to recompile the environment and optimize it for the new machine  
`spack concretize`  
`spack install`

# Most existing tools do not support combinatorial versioning

- Traditional binary package managers
  - RPM, yum, APT, yast, etc.
  - Designed to manage a single stack.
  - Install *one* version of each package in a single prefix (/usr).
  - Seamless upgrades to a *stable, well tested* stack
- Port systems
  - BSD Ports, portage, Macports, Homebrew, Gentoo, etc.
  - Minimal support for builds parameterized by compilers, dependency versions.
- Virtual Machines and Linux Containers (Docker)
  - Containers allow users to build environments for different applications.
  - Does not solve the build problem (someone has to build the image)
  - Performance, security, and upgrade issues prevent widespread HPC deployment.



# Modern scientific codes rely on icebergs of dependency libraries



Todd Gamblin, Gregory Becker, Alec Scott. Managing HPC Software Complexity with Spack. HPCIC Tutorials 2024, Livermore, California. July 22, 2024.

<https://spack-tutorial.readthedocs.io/en/latest/>

# Spack is not the only HPC/AI/data science package manager out there



## 1. “Functional” Package Managers

- Nix
- Guix

<https://nixos.org>  
<https://hpc.guix.info>



## 2. Build-from-source Package Managers

- Homebrew, LinuxBrew
- MacPorts
- Gentoo

<https://brew.sh>  
<https://www.macports.org>  
<https://gentoo.org>

## Other tools in the HPC Space:

### ■ Easybuild

- An installation tool for HPC
- Focused on HPC system administrators – different package model from Spack
- Relies on a fixed software stack – harder to tweak recipes for experimentation

<https://easybuild.io>

### ■ Conda / Mamba / Pixi

- Very popular binary package ecosystem for data science
- Not targeted at HPC; generally has unoptimized binaries

<https://conda.io>  
<https://mamba.readthedocs.io>  
<https://prefix.dev>



# How to learn more about Spack

- Main documentation: <https://spack.readthedocs.io/en/latest/>
- Spack tutorial
  - <https://spack-tutorial.readthedocs.io/en/latest/>
  - Video part 1: <https://www.youtube.com/watch?v=SShzurXZr4w> (3 hours)
  - Video part 2: <https://www.youtube.com/watch?v=fhijfzbVCH8> (3 hours)

- Part 1



- Part 2





# Spack in 1 slide

I want to build the OpenBLAS linear algebra library with GCC 12.3.0 and OpenMP multithreading enabled.

```
$ source spack/share/spack/setup-env.sh
```

Initialize Spack

```
$ spack spec -I openblas threads=openmp %gcc@12.3.0  
[some output removed]
```

Ask Spack how it would install OpenBLAS this way. "Dry run"

```
$ spack install openblas threads=openmp %gcc@12.3.0  
[some output removed]
```

Install!

```
==> openblas: Successfully installed openblas-0.3.26-r7xl6gpe2esp3qwpmj71qjmo2mwdt62  
Stage: 0.77s. Edit: 0.00s. Build: 4m 16.69s. Install: 1.49s. Post-install: 0.09s. Total: 4m 19.13s  
[+] /home/pla/spack/opt/spack/linux-ubuntu22.04-zen2/gcc-12.3.0/openblas-0.3.26-r7xl6gpe2esp3qwpmj71qjmo2mwdt62
```

```
$ ls /home/pla/spack/opt/spack/linux-ubuntu22.04-zen2/gcc-12.3.0/openblas-0.3.26-r7xl6gpe2esp3qwpmj71qjmo2mwdt62/lib  
cmake libopenblas.a libopenblas.so libopenblas.so.0 libopenblas-r0.3.26.a libopenblas-r0.3.26.so pkgconfig
```

```
$ spack load openblas /r7xl6gp
```

Load the software into your shell environment (PATH, CMAKE\_PREFIX\_PATH etc)

```
$ echo $CMAKE_PREFIX_PATH
```

```
/home/pla/spack/opt/spack/linux-ubuntu22.04-zen2/gcc-12.3.0/openblas-0.3.26-r7xl6gpe2esp3qwpmj71qjmo2mwdt62:/home/pla/spack/opt/spack/linux-ubuntu22.04-zen2/gcc-12.3.0/gcc-runtime-12.3.0-25tpakyzumeqrsmxzqcnsd4uifyzvrf
```

A CMake build would now be able to pick up OpenBLAS

# Spack provides a *spec* syntax to describe customized package configurations

```
$ spack install mpileaks                unconstrained
$ spack install mpileaks@3.3            @ custom version
$ spack install mpileaks@3.3 %gcc@4.7.3 % custom compiler
$ spack install mpileaks@3.3 %gcc@4.7.3 +threads +/- build option
$ spack install mpileaks@3.3 cppflags="-O3 -g3" set compiler flags
$ spack install mpileaks@3.3 target=cascadelake set target microarchitecture
$ spack install mpileaks@3.3 ^mpich@3.2 %gcc@4.9.3 ^ dependency constraints
```

- Each expression is a *spec* for a particular configuration
  - Each clause adds a constraint to the spec
  - Constraints are optional – specify only what you need.
  - Customize install on the command line!
- Spec syntax is recursive
  - Full control over the combinatorial build space

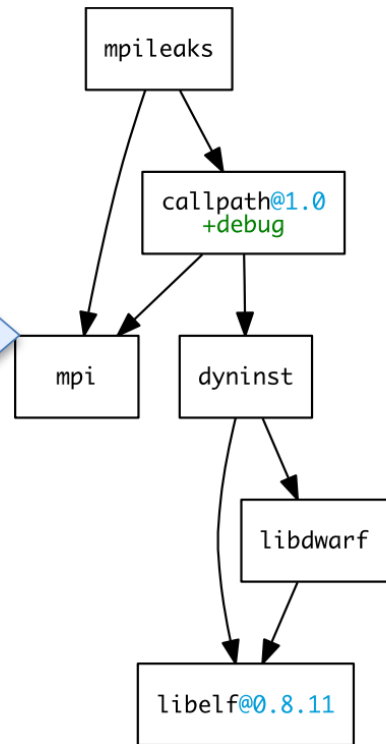
# Concretization fills in missing configuration details when the user is not explicit.

mpileaks ^callpath@1.0+debug ^libelf@0.8.11

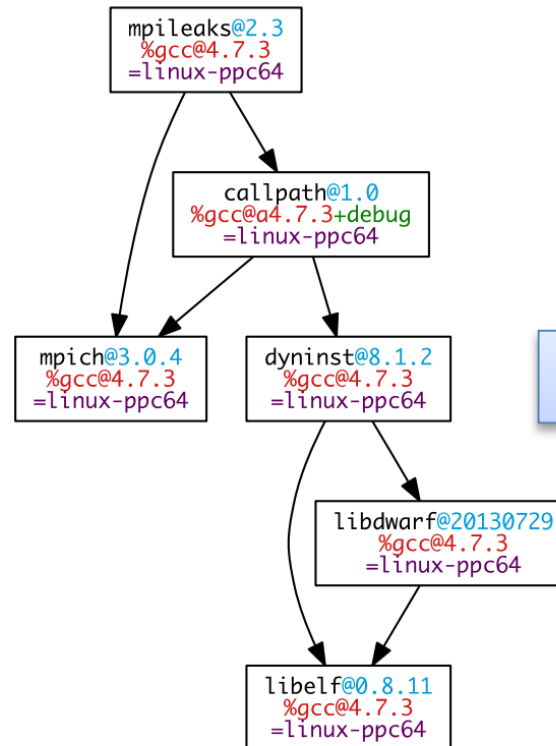
User input: *abstract* spec with some constraints

spec.json

Normalize



Concretize



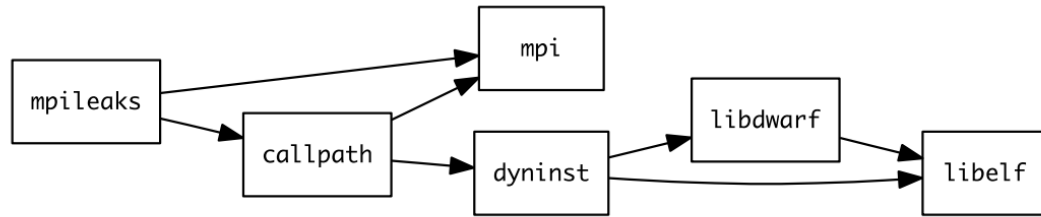
Store

```
{
  "spec": {
    "_meta": {
      "version": 4
    },
    "nodes": [
      {
        "name": "mpileaks",
        "version": "1.0",
        "arch": {
          "platform": "linux",
          "platform_os": "rhel8",
          "target": {
            "name": "cascadelake"
          }
        },
        "compiler": {
          "name": "gcc",
          "version": "10.3.1"
        },
        "namespace": "builtin",
        "parameters": {
          "build_system": "autotools",
          "stackstart": "0",
          "hash": "07awlh5q6wccrraon4yd2mfmkdtvvnxe"
        },
        "dependencies": [
          {
            "name": "adept-utils",
            "hash": "r3s7ywbhvtixgc3bknpqovl3dn2adce2",
            "parameters": {
              "deptypes": [
                "build",
                "link"
              ],
              "virtals": []
            }
          }
        ]
      }
    ]
  },
}
```


Detailed provenance stored with installed package

# Spack handles combinatorial software complexity

## Dependency DAG



## Installation Layout



```
opt
├── spack
│   ├── linux-rhel7-skylake
│   │   ├── gcc-8.3.0
│   │   │   ├── mpileaks-1.0-hc4sm4vuzpm4znmvrfzri4ow2mkphe2e
│   │   │   ├── callpath-1.0.4-daqqpssxb6qbfzrtsezkmhus3xoflbsy
│   │   │   ├── openmpi-4.1.4-u64v26igxvyn23hysmkdfums6tgjv5r
│   │   │   ├── dyninst-12.1.0-u64v26igxvyn23hysmkdfums6tgjv5r
│   │   │   ├── libdwarf-20180129-u5eawkvaoc7vonabe6nndkcfwuv233cj
│   │   │   └── libelf-0.8.13-x46q4wm46ay4pltrijbgizxjrhbaka6
```

- Each unique dependency graph is a unique **configuration**.
- Each configuration in a unique directory.
  - Multiple configurations of the same package can coexist.
- **Hash** of entire directed acyclic graph (DAG) is appended to each prefix.
- Installed packages automatically find dependencies
  - Spack embeds RPATHs in binaries.
  - No need to use modules or set `LD_LIBRARY_PATH`
  - Things work *the way you built them*

# We can configure Spack to build with external software

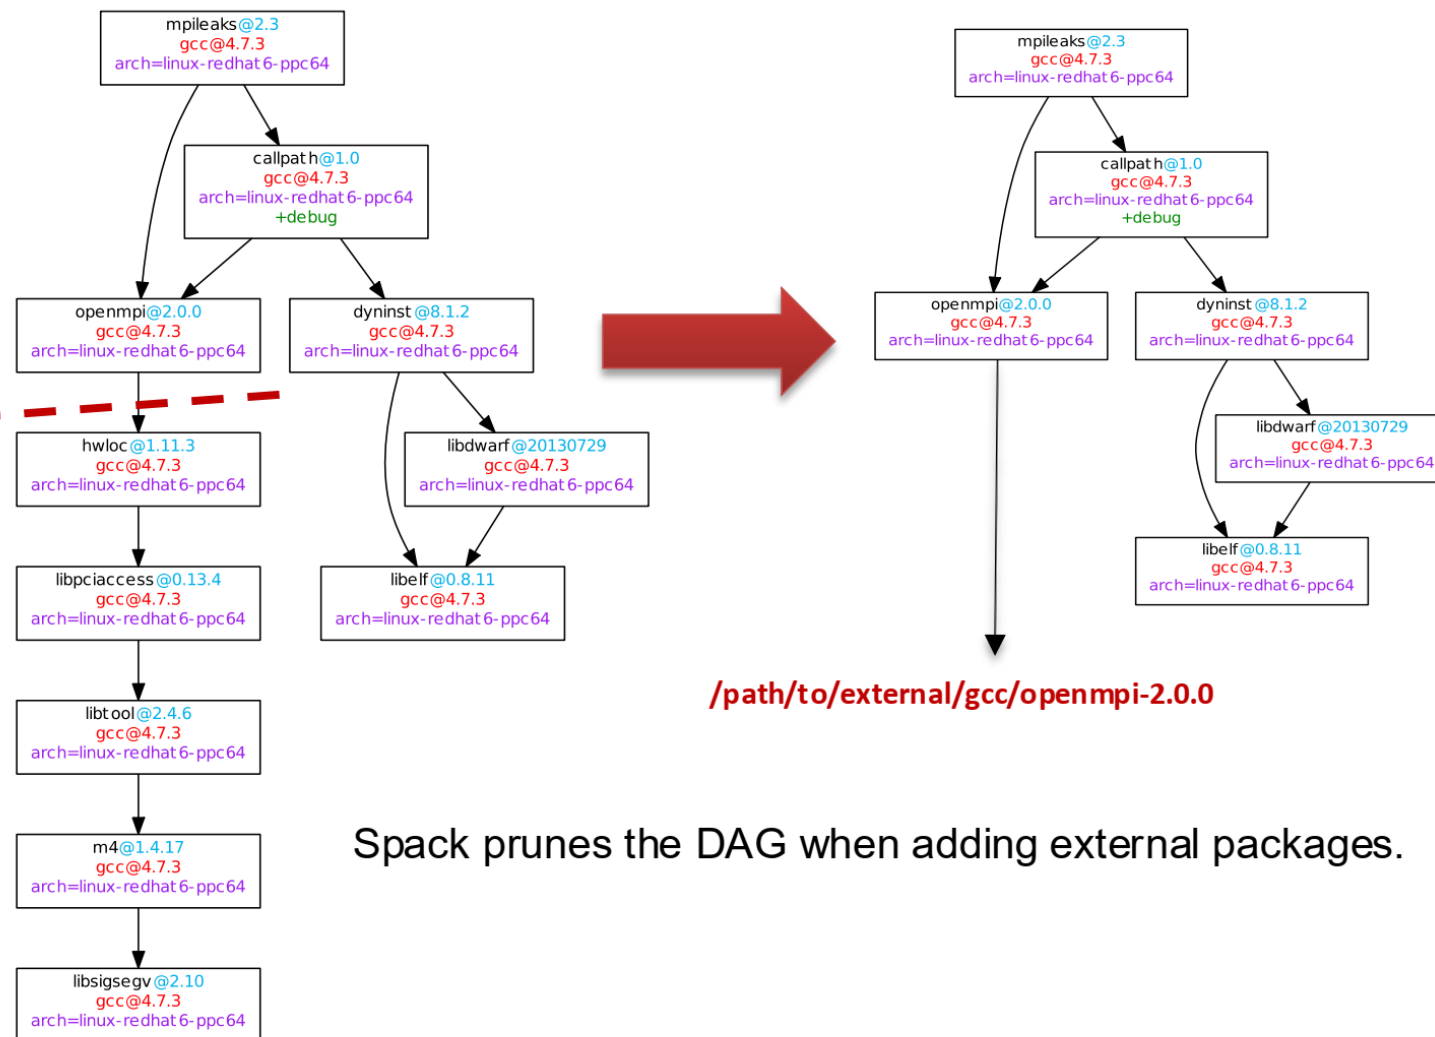
**mpileaks ^callpath@1.0+debug  
^openmpi ^libelf@0.8.11**

Build using the system's OpenMPI

## packages.yaml

```
packages:  
mpi:  
  buildable: False  
  paths:  
    openmpi@2.0.0 %gcc@4.7.3 arch=linux-rhel6-ppc64:  
      /path/to/external/gcc/openmpi-2.0.0  
    openmpi@1.10.3 %gcc@4.7.3 arch=linux-rhel6-ppc64:  
      /path/to/external/gcc/openmpi-1.10.3  
    ...
```

Users register external packages in a configuration file (more on these later).



Spack prunes the DAG when adding external packages.